Communication frameworks for enterprise architecture transformation

The most effective enterprise communication operates on a pyramid of five interlocking disciplines: executive briefing structures, stakeholder influence management, cross-functional negotiation, progress reporting, and architecture decision communication. In financial services transformation programs, these disciplines are not optional — they are the connective tissue between strategy and execution. This report catalogs **42 frameworks** across these five domains plus canvas-based strategy tools, maps them to evaluation rubrics, identifies existing open-source schemas, and provides the architectural foundation for a composable JSON spec kit designed for agentic AI integration.

The research reveals a critical gap: while individual frameworks are well-documented, **no open standard or JSON schema exists for business architecture communication artifacts** — no formal schemas for RACI matrices, OKRs, stakeholder maps, risk registers, or canvas frameworks. The Structurizr JSON schema (C4 model) GitHub and ArchiMate XML Exchange Format are the closest precedents. This gap represents the core opportunity for the "claude-power-pack" project.

Executive briefings demand answer-first structures

Seven frameworks dominate executive communication in financial services, all sharing one principle: **lead with the conclusion, then support it**.

The Pyramid Principle (Barbara Minto, 1970s/McKinsey) remains the gold standard. Substack Substack Its three-tier structure — governing thought at the top, **3 supporting arguments** in the middle, detailed evidence at the base — forces answer-first communication. Three rules govern construction: ideas at any level summarize those below, ideas in each group are logically the same kind, and ideas within groups follow logical order (time, structure, or degree). StrategyU Minto's invention of the **MECE principle** (Mutually Exclusive, Collectively Exhaustive) provides the analytical backbone for structuring any complex argument without gaps or overlaps. Wikipedia

SCQA (Situation-Complication-Question-Answer), also from Minto, functions as the narrative wrapper around the pyramid. Fpa-trends It establishes baseline context, introduces the tension demanding attention, focuses the audience on the decision point, then delivers the recommendation. Substack Analytic-storytelling BCG uses a three-step variant (SCR — Situation-Complication-Resolution). Slide Science In financial services, SCQA is particularly effective for framing transformation updates: *"Our core banking platform handles branch transactions [S]. Digital competitors capture 30% of new accounts [C]. How do we modernize while maintaining regulatory compliance? [Q]. A phased API-first migration over 18 months [A]."*

BLUF (Bottom Line Up Front), from U.S. military doctrine, strips communication to its essential judgment. Wikipedia Every fact must pass the "So What" test. Willis Towers Watson research shows companies with effective communication practices produce **47% higher shareholder returns**. The Persimmon Group Amazon's **6-Pager Memo** (Jeff Bezos, 2004) takes the opposite approach to brevity — banning PowerPoint in favor of

narrative prose (Six Pager Memo) across seven sections (Introduction, Goals, Tenets, State of Business, Lessons Learned, Strategic Priorities, Appendix). (LinkedIn) Memos are read silently for 20–25 minutes at meeting start, (Setonhill) forcing deeper analytical clarity than slides permit.

McKinsey's storyline structure layers on top: the GAME framework (Goal → Audience → Message → Expression), (Steveseager) action titles that tell the full story when read sequentially, and the governing thought principle — every communication has ONE key message. (Medium) McKinsey's internal research found **25% productivity gains** possible through structured communication. (McKinsey)

Framework	Creator	Core Principle	Best For	Format
Pyramid Principle	Minto (1970s)	Answer-first, MECE grouping	Board presentations, strategy papers	Any
SCQA	Minto	Narrative tension → resolution	Framing steering committee discussions	Presentations, emails
BLUF	U.S. Military	Lead with judgment	Status updates, risk alerts, decision memos	Written
McKinsey Storyline	McKinsey	Governing thought + GAME	Due diligence, M&A communications	Slides, reports
Amazon 6-Pager	Bezos (2004)	Deep narrative prose	Major strategic decisions	Written memo
Monroe's Motivated Sequence	Monroe (1930s)	Attention→Need→Satisfaction→Visualization→Action	Investment proposals, change pitches	Presentations
Executive 1-Pager	General practice	Decision-required + 3–5 key points	Gate reviews, quick approvals	Written

Monroe's Motivated Sequence deserves special mention for transformation business cases: its five-step psychological arc (Attention → Need → Satisfaction → Visualization → Action) mirrors how executives process investment decisions, moving from awareness through to commitment.

Stakeholder alignment requires layered influence management

Effective stakeholder management operates across three layers: **mapping** (who matters), **understanding** (their disposition), and **moving** (toward alignment).

For mapping, **Mendelow's Power-Interest Matrix** (1991) remains the most widely used tool Toolshero — a 2×2 grid plotting stakeholders on power and interest, Professional Academy producing four engagement strategies: Manage Closely (high power, high interest), Keep Satisfied (high power, low interest), Keep Informed (low power, high interest), and Monitor (low power, low interest). Oxford College of Marketing For more nuanced analysis, **Mitchell's Stakeholder Salience Model** (1997) adds legitimacy and urgency as dimensions, LinkedIn creating seven stakeholder categories from Dormant (power only) to Definitive (all three attributes). Wikipedia **The Stakeholder Circle** (Dr. Lynda Bourne, 2002) adds lifecycle awareness — stakeholder importance shifts across project phases. Passionate About OSS

For role clarity, **RACI** (Responsible, Accountable, Consulted, Informed) Openpracticelibrary is table stakes, Wikipedia but financial services programs benefit from pairing it with decision-focused variants. **DACI** (Driver, Approver, Contributor, Informed) maps decision authority. Kumospace Bawiki Bain's **RAPID** (Recommend, Agree, Perform, Input, Decide) specifically addresses who has the final "D" Umbrex — critical in matrix organizations where accountability is diffuse. In compliance-heavy environments, **RACI-VS** adds Verifier and Signatory roles. Kumospace Perfony

For moving stakeholders toward alignment, three frameworks stand out. **Kotter's 8-Step Change Management Model** (1995) provides the macro communication arc — from creating urgency through building coalitions, communicating vision, enabling action, generating quick wins, and institutionalizing change. Prosci Agile Academy **Prosci's ADKAR** (Awareness, Desire, Knowledge, Ability, Reinforcement) focuses on individual adoption, diagnosing exactly where each stakeholder is "stuck." Prosci Myriam Tisler Organizations using ADKAR are **73% more likely** to achieve change goals. Change-management-hub **Cialdini's 7 Principles of Persuasion** (1984/2016) — Reciprocity, Commitment, Social Proof, Authority, Liking, Scarcity, Unity CXL — provide tactical influence levers applicable to every stakeholder interaction.

For group decision-making, **Gradients of Agreement** (Sam Kaner, 1987) Wikipedia replaces binary yes/no with an 8-point continuum from "whole-hearted endorsement" to "veto." Integration and Implementation… This surfaces nuanced positions Integration and Implementation… and prevents the common pattern in steering committees where silence is mistaken for consensus. Integration and Implementation…

Cross-functional negotiation follows principled frameworks

Enterprise transformation programs live or die at the boundaries between business units. Three negotiation and conflict frameworks are essential.

Interest-based negotiation (Fisher & Ury, *Getting to Yes*, 1981) provides four principles: separate people from problems, focus on interests not positions, generate options for mutual gain, and insist on objective criteria.

Beyond Intractability +2 The companion concept of **BATNA** (Best Alternative to a Negotiated Agreement) establishes each party's walk-away benchmark Pressbooks Wikipedia — critical when Platform and Product teams disagree on roadmap sequencing. A team's BATNA might be building a workaround, using a third-party service, or escalating to leadership.

The **Thomas-Kilmann Conflict Mode Instrument** (1974) maps five modes along assertiveness and cooperativeness axes: Competing (win-lose), Collaborating (win-win), Compromising (split the difference), Avoiding (sidestep), and Accommodating (yield). The key insight is **no mode is inherently best** — collaborating suits critical architecture decisions, while compromising works for lower-priority scheduling disputes.

For dependency management at scale, **SAFe's PI Planning** provides the ceremony — a 2-day event where all Agile teams align on shared objectives and identify dependencies Easy Agile using the **Program Board** (teams × iterations, with red/gray strings showing dependency risk). **SLAs and OLAs** (formalized through ITIL) create formal internal agreements between teams, defining response times, escalation paths, and performance metrics. The critical distinction: SLAs are customer-facing commitments; OLAs are internal agreements specifying how teams collaborate to meet those commitments. Infraon Blogs Invgate

Dependency mapping itself follows four steps: identify all cross-team dependencies, plot them on an N×N matrix, visualize on a timeline board with color coding, and classify by type (technical, resource, process, external) and risk level (hard/blocking vs. soft/informational). Too many red strings on a program board signals the need for organizational redesign, not just better coordination.

Progress reporting combines quantitative rigor with visual clarity

Six reporting frameworks form a complementary stack, from tactical to strategic.

RAG status (Red/Amber/Green) provides instant visual comprehension BestOutcome but suffers from **"watermelon reporting"** — projects appearing green on the outside while red inside. Mitigation requires tying RAG to objective thresholds (e.g., Green = within 5% budget variance, Amber = 5–10%, Red = >10%) rather than subjective judgment. **RAID logs** (Risks, Assumptions, Issues, Dependencies) capture multiple dimensions of project health in one document, ProjectManager +3 reviewed weekly and escalated at monthly governance meetings. Asana

Earned Value Management delivers the most objective assessment through three core metrics: **CPI** (Cost Performance Index = EV/AC), **SPI** (Schedule Performance Index = EV/PV), and **EAC** (Estimate at Completion). 6 Sigma Research shows CPI at **20% project completion** reliably predicts final cost. PMI For executive communication, trend lines showing CPI/SPI over time are more valuable than point-in-time snapshots.

OKRs (Andy Grove/Intel, 1970s; popularized at Google by John Doerr, 1999) AltexSoft align organizations around outcomes rather than activities. Wikipedia The structure — 3–5 qualitative Objectives with 2–5 measurable Key Results each, scored 0.0–1.0 on a quarterly cadence IBM — drives focus and transparency.

The Balanced Scorecard (Kaplan & Norton, 1992) provides strategic context through four perspectives: Financial, Customer, Internal Process, and Learning & Growth, Harvard Business School connected by cause-and-effect strategy maps. Harvard Business School

Stephen Few's **dashboard design principles** ensure all reporting is cognitively effective: Amazon fit on a single screen, ResearchGate maximize data-ink ratio, use bullet graphs for performance against targets, sparklines for inline trends, Smart Cville and always provide context (targets, historical comparisons, benchmarks). His core maxim: **a number without context is meaningless**. Amazon

---

Architecture communication bridges technical depth and business clarity

Eight frameworks help architects communicate decisions, trade-offs, and integration approaches to diverse audiences.

Architecture Decision Records (Michael Nygard, 2011) are the most widely adopted lightweight format: each ADR captures Title, Status, Context, Decision, and Consequences in a short text file stored alongside code. Cognitect AWS The **MADR variant** adds explicit option analysis with pros/cons. Architectural Decision Records **Y-Statements** compress decisions to a single sentence: *"In the context of [situation], facing [concern], we decided for [option] to achieve [quality], accepting [downside]."* ThoughtWorks placed ADRs in its "Adopt" category. Red Hat

The **C4 Model** (Simon Brown, 2006–2011) provides a "Google Maps zoom" approach to architecture visualization: Level 1 (Context) for anyone, Level 2 (Containers) for mostly technical audiences, Level 3 (Components) for developers and architects, Level 4 (Code) for implementers. InfoQ +3 **Capability Mapping** gives leadership a stable, technology-neutral vocabulary — a hierarchical taxonomy of what a business does (not how), typically 3 levels deep, enabling heat-map overlays for maturity, investment, and risk. The Open Group Orbus Software

Value Stream Mapping (Toyota/Rother & Shook, 1998) visualizes end-to-end value delivery from order to fulfillment, identifying waste and handoff friction. Lean Enterprise Institute In financial services, it maps processes like client onboarding or claims processing. **TOGAF** provides the comprehensive enterprise architecture framework with viewpoints tailored to different stakeholder categories, The Open Group The Open Group while **ArchiMate** (~50 concepts across Business, Application, and Technology layers) Wikipedia offers the modeling language. Wikipedia The **Zachman Framework** (1987) provides the completeness check — a 6×6 matrix of interrogatives (What, How, Where, Who, When, Why) against stakeholder perspectives. LeanIX ScienceDirect

For evaluating architectural trade-offs, **ATAM** (SEI/CMU, 1998) is the gold standard: Software Engineering Institute a 3–4 day structured evaluation involving Software Engineering Institute 15–30 participants that produces a utility tree of quality attributes, sensitivity points, and trade-off points — exposing where improving one quality attribute necessarily degrades another. Software Engineering Institute Software Engineering Institute

---

Canvas frameworks map strategy on a single page

Seven canvas-based tools provide visual strategy communication, each serving a distinct purpose.

The **Business Model Canvas** (Osterwalder, 2010) remains dominant with its 9 building blocks Ixdf ThePowerMBA spanning value creation, delivery, and capture. ResearchGate Ixdf Its companion, the **Value Proposition Canvas**, zooms into customer-product fit through 6 segments: Customer Jobs, Pains, and Gains mapped against Products & Services, Pain Relievers, and Gain Creators. Ixdf Linden Innovation The **Lean Canvas** (Ash Maurya, 2012) replaces BMC's Key Partners, Activities, Resources, and Relationships with Problem, Solution, Key Metrics, and Unfair Advantage Bmtoolbox GitHub — optimized for startups testing hypotheses rather than documenting established models. Miro

Wardley Maps (Simon Wardley, 2005) are uniquely powerful for financial services strategy. They plot components on two axes: **value chain** (visibility to user) on the Y-axis and **evolution** (Genesis → Custom → Product → Commodity) on the X-axis. Wikipedia This reveals which banking technology components are ripe for commoditization, where custom investment creates competitive advantage, and where disruption risk lies. The **Strategy Canvas** (Kim & Mauborgne, Blue Ocean, 2005) complements by mapping value curves across 8–12 competing factors, exposing differentiation opportunities via the Eliminate-Reduce-Raise-Create framework. Umbrex

The **Platform Canvas** (Allweins/Proesch/Ladd, 2019) with 12 components Wikipedia addresses what BMC cannot: multi-sided dynamics, network effects, cross-side pricing, and governance FourWeekMBA — critical for open banking platforms and payment ecosystems. **Customer Journey Maps** complete the picture by visualizing end-to-end experience stages, touchpoints, emotions, and pain points. Business Models Inc.

Agentic AI is reshaping project management at enterprise scale

The landscape of AI-augmented project management has moved from experimental to operational. Gartner projects **40% of enterprise applications** will integrate task-specific AI agents by end of 2026, with **15% of day-to-day work decisions** made autonomously by 2028. The AI project management market reached **$3.08B in 2024** and is projected to hit $7.4B by 2029.

Leading platforms include **Asana AI** (AI Teammates receiving task assignments like human team members, UC Today AI Risk Reports generating weekly risk assessments), Asana **Microsoft Planner + Copilot** (Project Manager Agent breaking goals into tasks and executing them), **Smartsheet** (Knowledge Graph for org-specific AI context, Smart Agents for continuous monitoring), UC Today and the open-source **Agentic Project Management framework** on GitHub (multi-agent coordination with Cursor, Claude Code, and GitHub Copilot integration). GitHub

Three emerging protocols matter for financial services: **Model Context Protocol (MCP)** from Anthropic for horizontal agent-to-tool connections, **Agent2Agent Protocol (A2A)** from Google for vertical agent orchestration, and **Agent Communication Protocol (ACP)** for offline collaboration. Deloitte Insights In banking

specifically, early adopters report **20% faster processing times** and **15% lower operational costs** (nCino research), while Moody's users consume **60% more research** while cutting task completion time by 30%.

However, Gartner warns that **over 40% of agentic AI projects will be canceled by end of 2027** due to escalating costs, unclear business value, or inadequate risk controls. Only 26% of companies can move beyond proof of concept to create value at scale (Capgemini). The governance imperative is clear: enterprises need **Guardian Agents** (Reviewers, Monitors, Protectors) overseeing other agents, and must adopt a "self-service with guardrails" model similar to Microsoft's internal approach. Microsoft

Architecture as Code is the CI/CD parallel for business architecture. medium Structurizr DSL enables version-controlled architecture models that generate diagrams automatically, integrate with CI/CD pipelines, and catch architectural drift. ADRs stored in Git enable pull-request-based architecture review. Cognitect +3 The shift from plan-build-run to continuous business change mirrors DevOps transformation Ardoq — enterprise architecture artifacts become living, versioned, continuously deployed assets.

Evaluation rubrics make communication quality measurable and AI-evaluable

Structuring communication criteria as machine-readable rubrics enables both human assessment and AI-powered quality gates. The research identifies a clear design pattern:

Each rubric contains **criteria** (dimensions to evaluate), **performance levels** (typically 4–5 gradations), **descriptors** (observable characteristics at each level), and **weights** (relative importance). DeepSeek AI Best practices call for 3–5 performance levels, 2+ distinct criteria, and descriptors that distinguish qualitative differences in observable, measurable terms. Teach Online

Eight scoring dimensions emerge as universal for project communication quality:

- **Clarity** — Unambiguous language and well-organized structure
- **Completeness** — All required sections populated with sufficient detail
- **Actionability** — Clear next steps, owners, and deadlines identified
- **Accuracy** — Factually correct and consistent with other artifacts
- **Audience appropriateness** — Tailored to stakeholder level (executive vs. technical)
- **Timeliness** — Delivered within required timeframes
- **Traceability** — Links to decisions, requirements, and dependencies
- **Risk transparency** — Honest assessment with mitigation plans

The **Microsoft 365 Communication Maturity Model** provides the most structured maturity assessment: Level 100 (Initial/ad hoc), Level 200 (Managed/repeatable), Level 300 (Defined/documented), Level 400 (Measured/high trust), Level 500 (Optimizing/automated). Microsoft Learn **Bloom's Taxonomy** structures evaluation depth: Remember (does the artifact contain required info?) → Understand (is it comprehensible?) →

Apply (does it correctly use frameworks?) → Analyze (does it identify trade-offs?) → Evaluate (does it make justified recommendations?) → Create (does it synthesize novel solutions?). Distancelearning

For AI evaluation, Amazon Nova's rubric-based LLM judge demonstrates dynamic rubric generation where criteria and weights are generated per-prompt context. AWS Snorkel AI emphasizes three rubric qualities for AI use: Scalability, Reliability (concrete enough for prompt-based scoring), and Alignment (with expert opinion). Snorkel AI Organizations with structured prompt engineering processes report **34% higher satisfaction** with AI implementations Profiletree and **76% fewer hallucination errors**. Profiletree

---

Existing open-source schemas reveal a major gap to fill

The most mature ecosystem is **ADR tooling**: joelparkerhenderson/architecture-decision-record provides comprehensive templates, GitHub GitHub adr/madr (v4.0.0) offers the most structured Markdown+YAML format, Architectural Decision Records and log4brains adds CLI tooling with CI/CD publishing. GitHub For enterprise architecture, Nasdanika-Models/togaf represents TOGAF 10 concepts as YAML/Markdown/Drawio — the closest existing repo to "enterprise architecture as code." GitHub

The **Structurizr JSON schema** (structurizr/json, Apache-2.0) is the only formal JSON schema for architecture modeling, github covering workspaces, systems, containers, components, deployment nodes, views, and documentation. GitHub The **ArchiMate Model Exchange File Format** (XML/XSD, mandatory for certified tools since 2018) is the only standards-body-backed interchange format for EA models. The Open Group +2

No JSON schema exists for: Business Model Canvas, Value Proposition Canvas, Lean Canvas, Wardley Maps, RACI matrices, OKR specifications, stakeholder maps, risk registers, status reports, RAID logs, or any project communication artifact. No OASIS, OMG, or ISO standard covers these artifacts in structured format. This represents the **core value proposition** of the claude-power-pack spec kit: defining composable, machine-readable schemas for the full spectrum of business architecture and project communication artifacts.

The recommended schema architecture for the spec kit should follow this hierarchy: **Categories** (the 5 communication contexts + canvas tools + agentic AI workflows) → **Frameworks** (individual named frameworks with metadata) → **Principles** (core rules and guidelines per framework) → **Evaluation Rubrics** (dimensions × levels × descriptors × weights) → **Composable Criteria** (reusable quality gates applicable across frameworks). Each framework entry should include: id, name, category, creator, core_structure (array of components), audience_levels, when_to_use, strengths, limitations, financial_services_context, evaluation_rubric (linked rubric ID), related_frameworks, and template (the actual artifact structure). Rubrics should follow the pattern: { criteria: [{ id, name, description, weight, levels: [{ score, label, descriptor }] }] }.

Conclusion

This research maps a complete framework landscape for enterprise transformation communication — from Minto's Pyramid Principle for board presentations to Structurizr's JSON schema for architecture-as-code. The

most actionable insight for the claude-power-pack project is that **the frameworks are well-established but the schemas don't exist**. No standard defines structured, machine-readable formats for the artifacts that enterprise programs produce daily: stakeholder maps, RACI matrices, status reports, risk registers, canvas analyses, or architecture decisions beyond ADRs. The project's unique contribution will be defining these schemas — enabling AI agents to generate, evaluate, and iterate on communication artifacts using rubric-based quality gates. The eight universal scoring dimensions (clarity, completeness, actionability, accuracy, audience appropriateness, timeliness, traceability, risk transparency), combined with Bloom's cognitive hierarchy TIPS and 5-level maturity models, provide the evaluation architecture. The agentic AI ecosystem (MCP, A2A, Guardian Agents) provides the runtime. What's missing is the structured specification layer in between — and that is precisely what this spec kit should deliver.